




UNIVERSITÀ
DI TORINO

Analisi e Visualizzazione delle Reti Complesse

NS22 - Traditional Machine Learning on Graphs


Prof. Rossano Schifanella




General Pipeline

(Supervised) Machine Learning Lifecycle

This feature, that feature. **Every single time!**



Different types of tasks

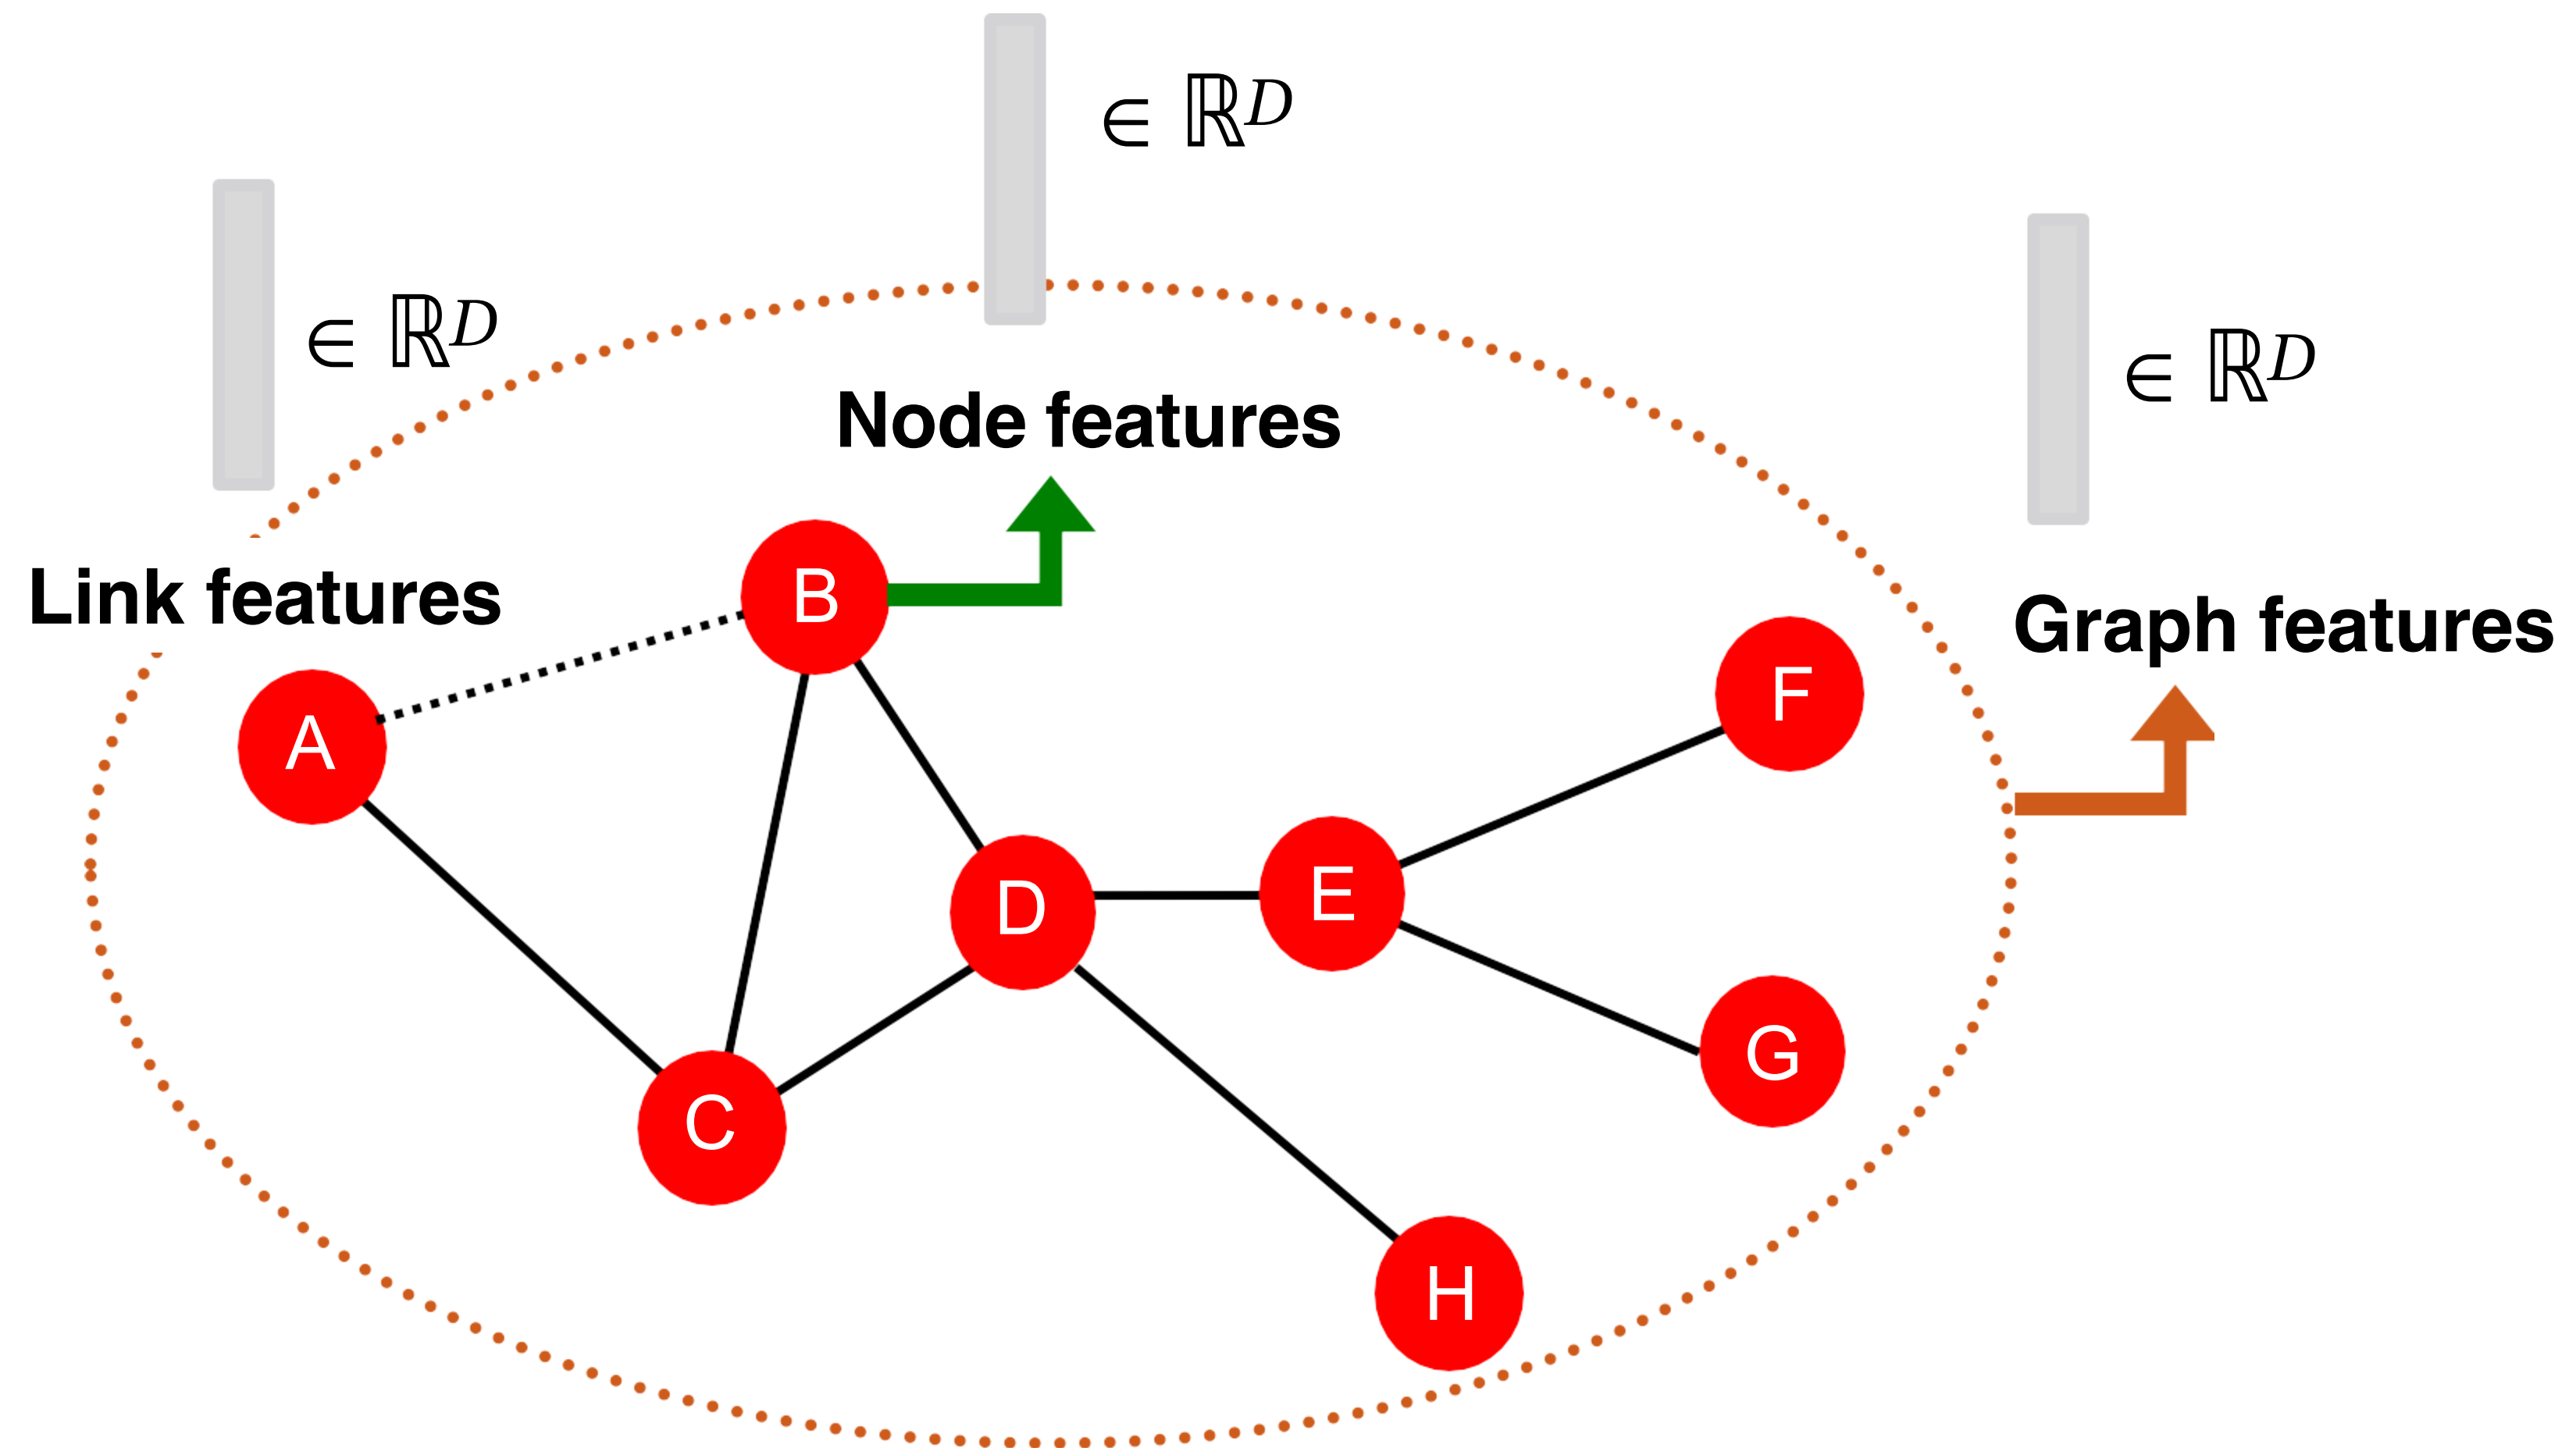


Classic Graph ML Tasks

- **Node classification:** predict a property of a node
 - Example: Categorize online users/items
- **Link prediction:** predict whether there are missing links between two nodes
 - Example: Knowledge graph completion
- **Graph classification:** categorize different graphs
 - Example: Molecule property prediction
- **Clustering:** detect if nodes form a community
 - Example: Social circle detection
- **Other tasks:**
 - **Graph generation:** Drug discovery
 - **Graph evolution:** Physical simulation

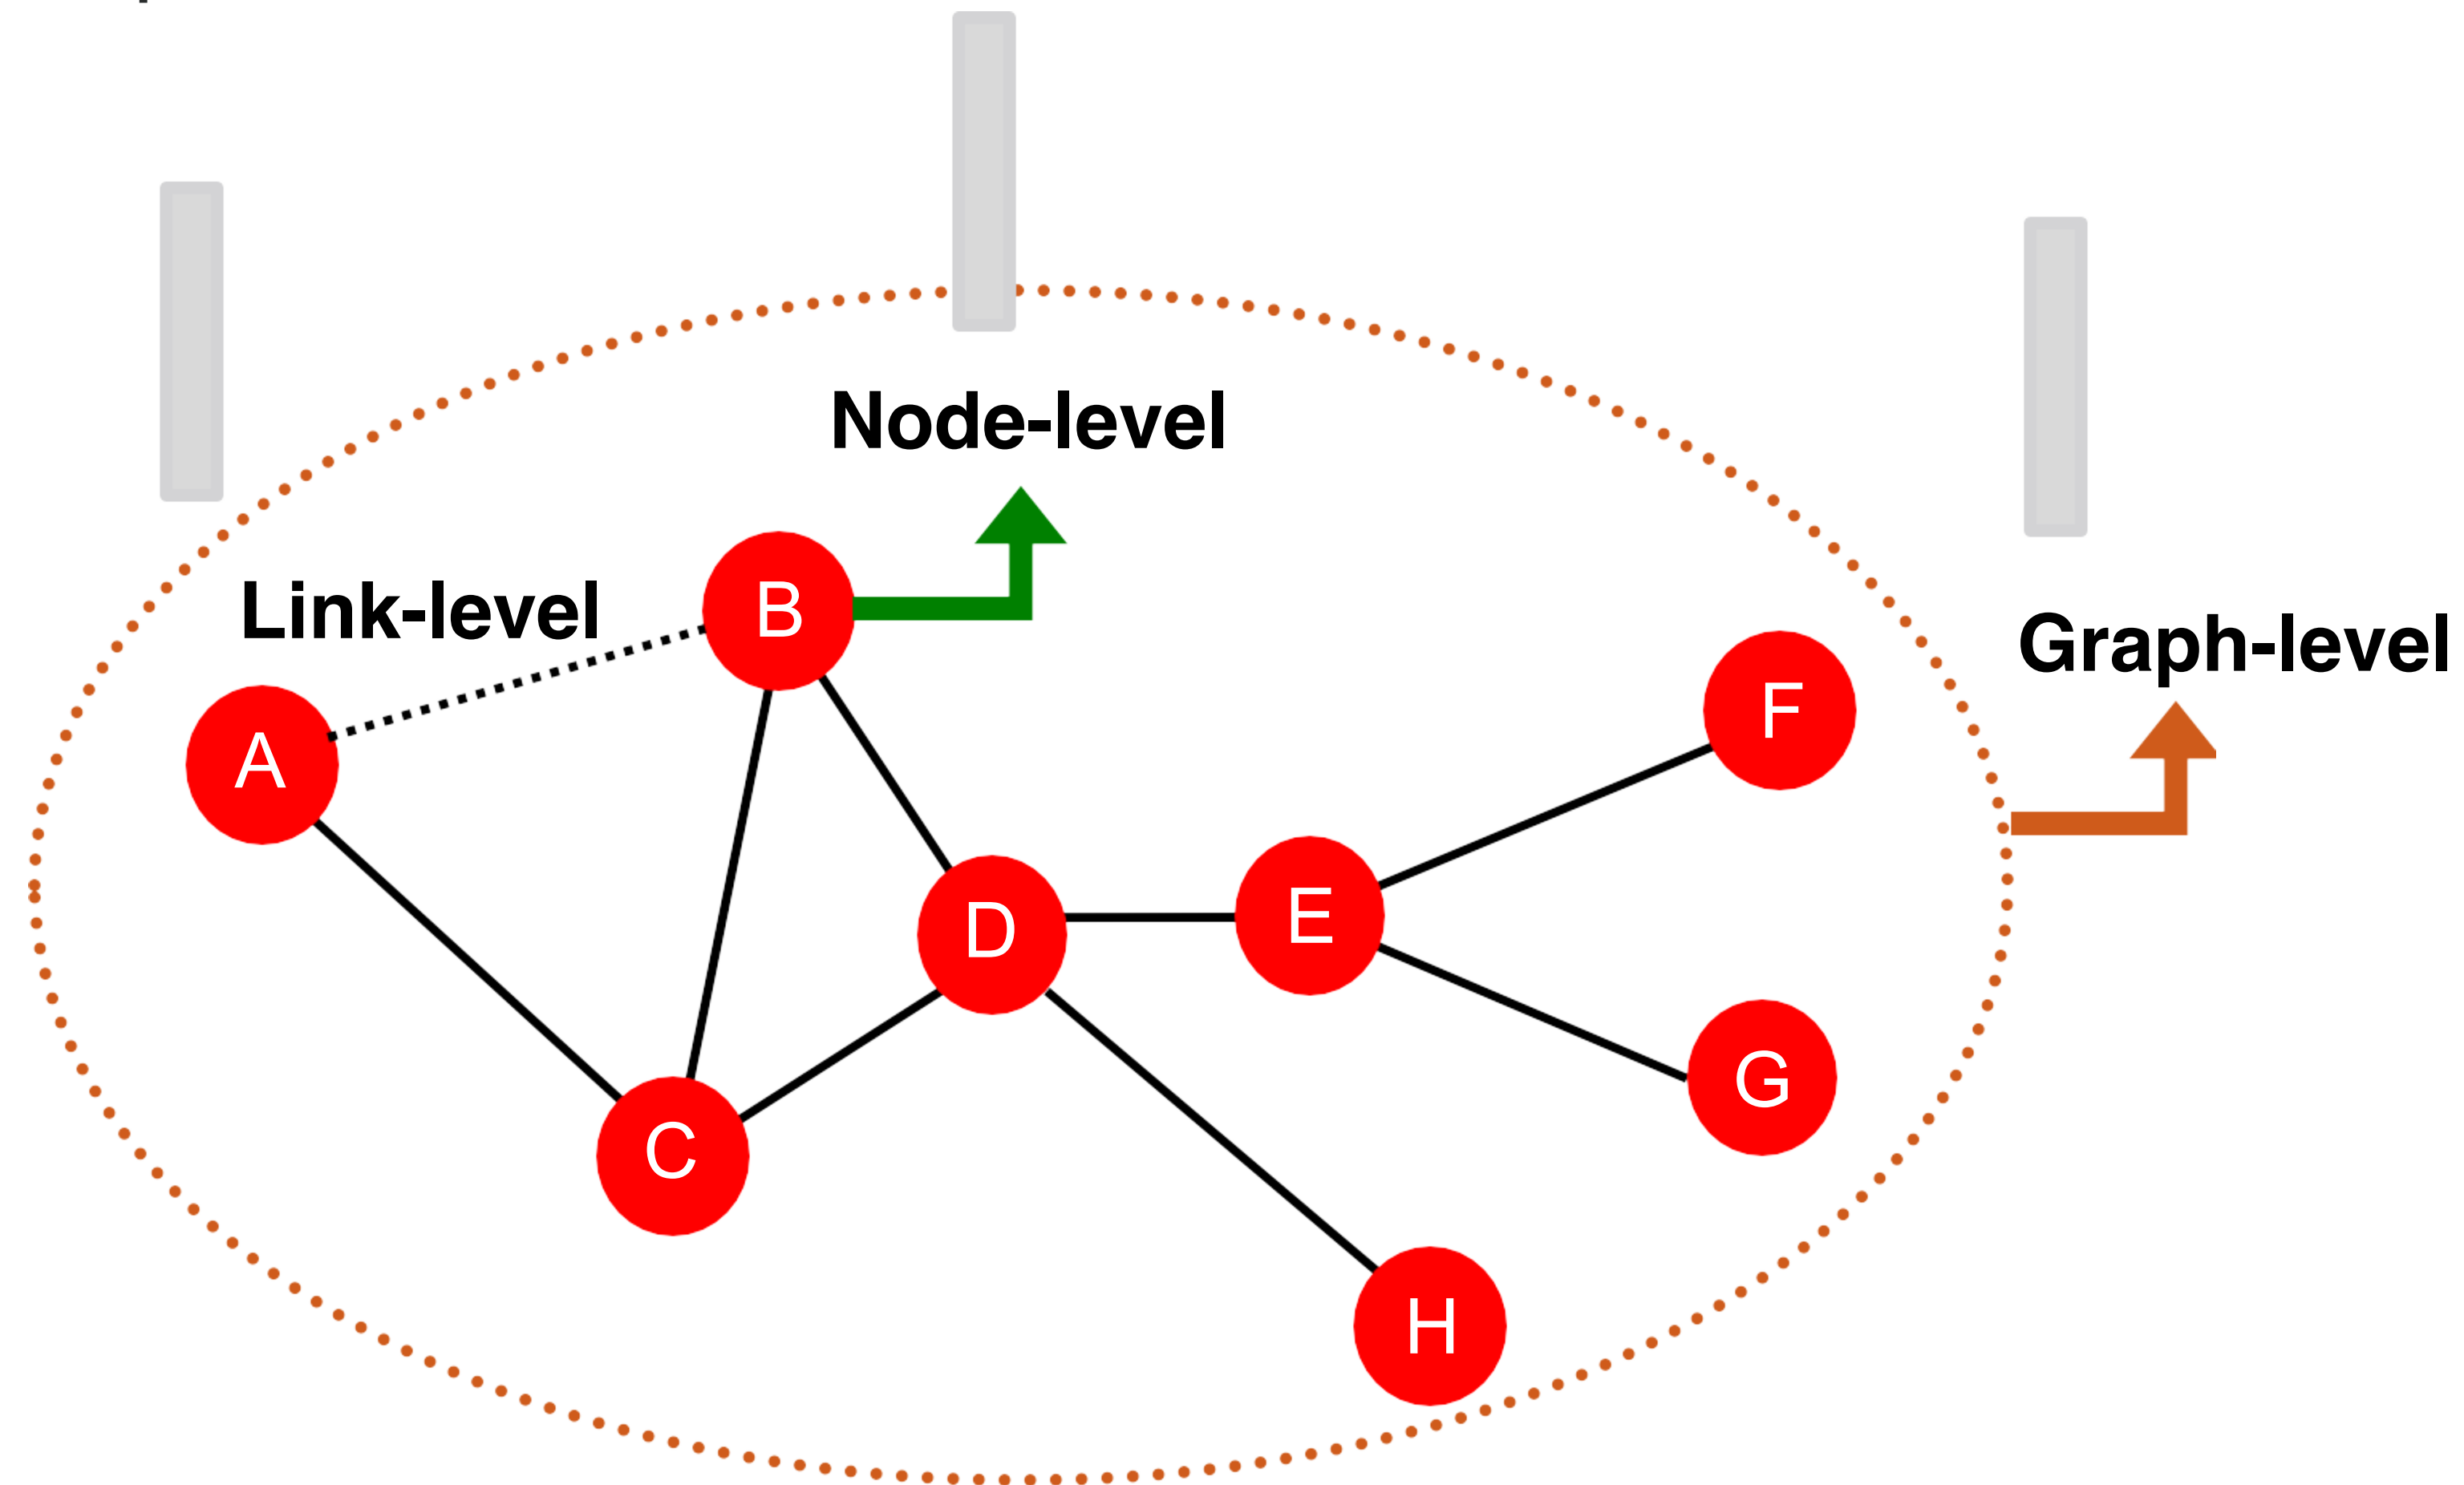
Traditional LM Pipeline

- Design features for nodes/links/graphs
- Obtain features for all training data



Machine Learning Tasks: Review

- **Node-level** prediction
- **Link-level** prediction
- **Graph-level** prediction





Traditional LM Pipeline

- **Train an ML model:**

- Logistic Regression
- Random forest
- Neural network, etc.

- **Apply the model:**

- Given a new node/link/graph
- Obtain its features
- Make a prediction



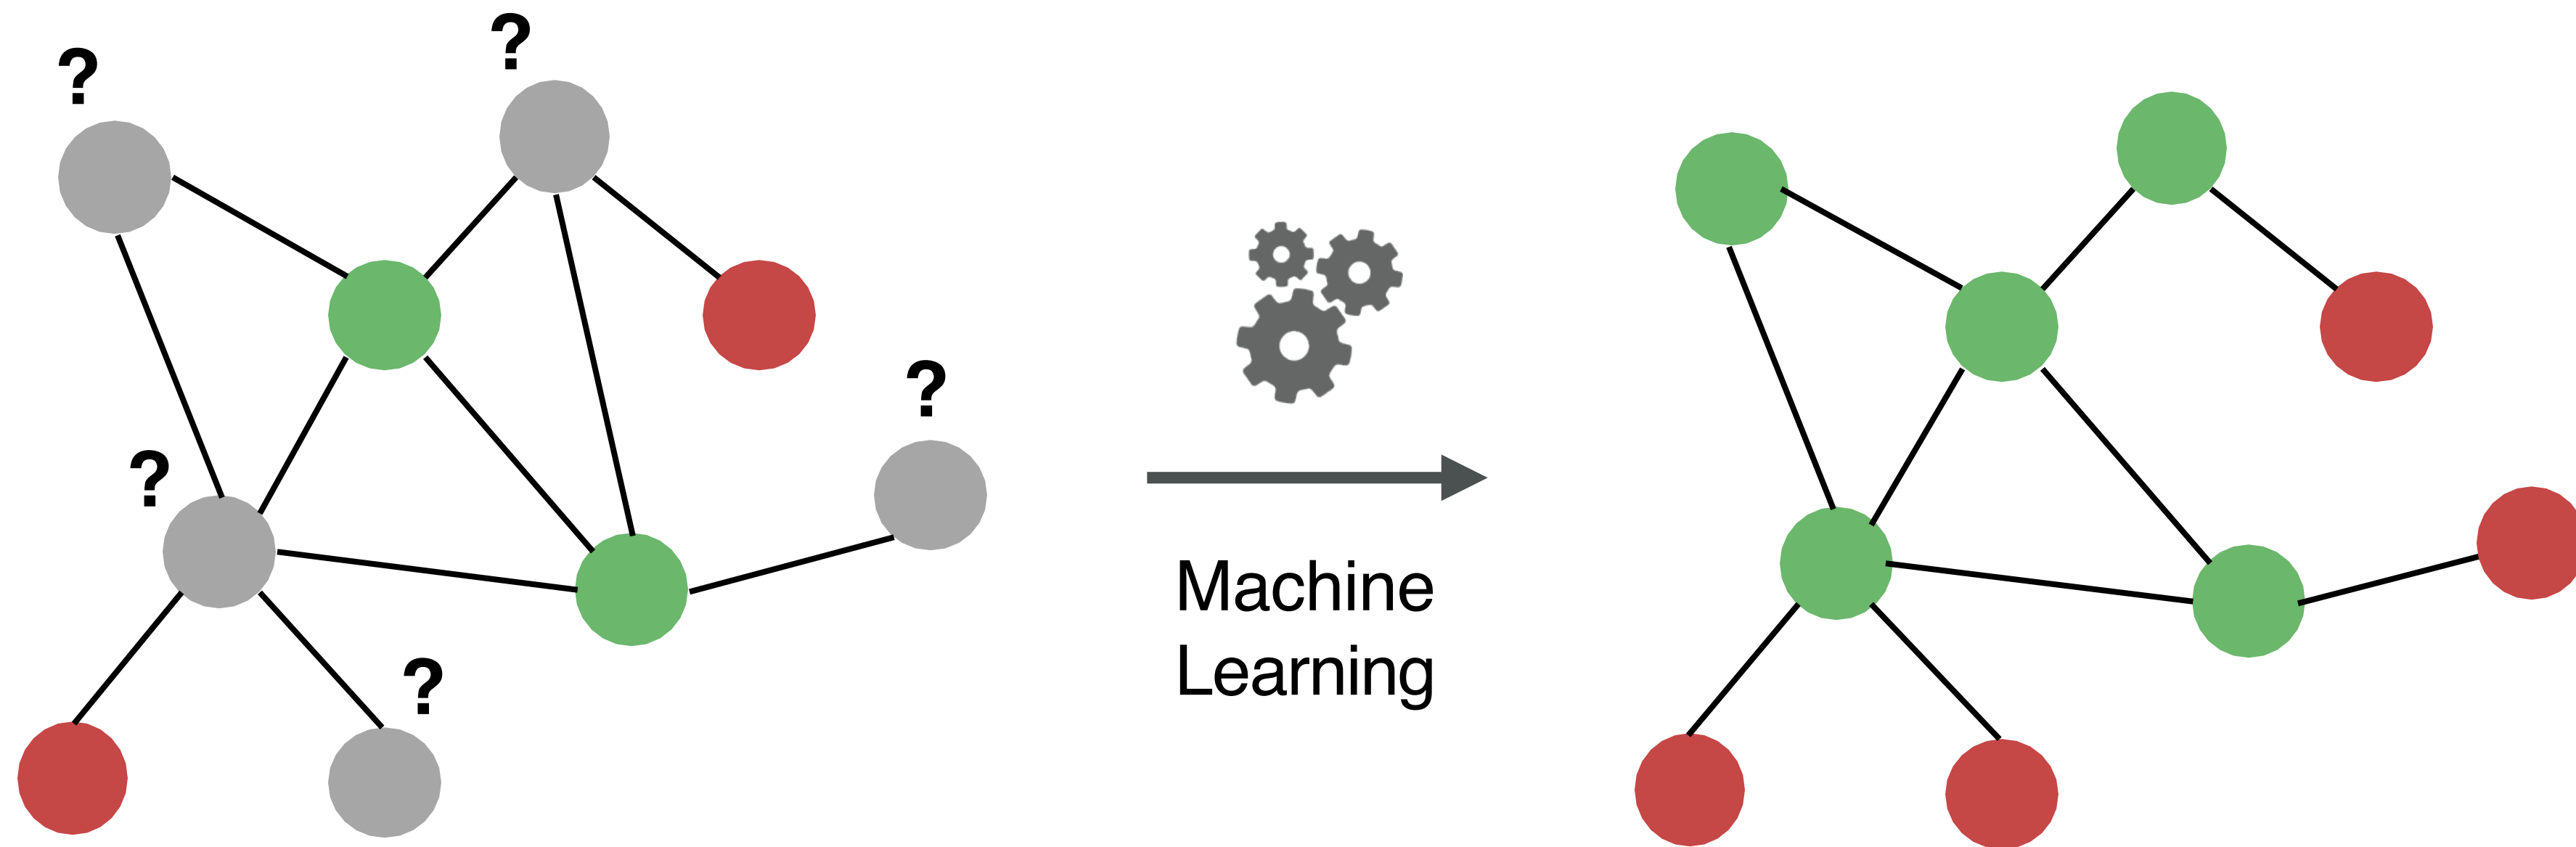
Features Design

- Using **effective features** x over graphs is the key to achieving good model performance.
- Traditional ML pipeline uses **hand-designed features**.
- In this lecture, we overview the traditional features for:
 - **Node-level** prediction
 - **Link-level** prediction
 - **Graph-level** prediction
- For simplicity, we focus on **undirected** graphs.

Machine Learning in Graphs

- **Goal:** Make predictions for a set of objects
- Design choices:
 - **Features:** d-dimensional vectors x
 - **Objects:** Nodes, edges, sets of nodes, entire graphs
 - **Objective function:**
 - What task are we aiming to solve?

Node-Level Tasks



Node classification

ML needs features.


Node-Level Overview

Goal

Characterize the structure and position of a node in the network:


- Node degree
- Node centrality
- Clustering coefficient
- **Graphlets**

Node features




Node Features: Graphlets

- Observation: Clustering coefficient counts the **#triangles** in the ego network.



$$e_v = 0.5$$



3 triangles (out of 6 node triplets)


- We can generalize the above by counting **#(pre-specified subgraphs, i.e., graphlets)**

Node Features: Graphlets

- **Goal:** Describe the network structure around node u
- **Graphlets are small subgraphs that describe the structure of node u 's network neighborhood.**



Analogy:

- Degree counts **#(edges)** that a node touches
- Clustering coefficient counts **#(of triangles)** that a node touches.
- **Graphlet Degree Vector (GDV):** Graphlet-base features for nodes
 - GDV counts **#(graphlets)** that a node touches




Node Features: Graphlets

- Considering graphlets of size 2-5 nodes we get:
 - A vector of 73 coordinates is a signature of a node that describes the topology of the node's neighborhood
- Graphlet degree vector provides a measure of a **node's local network topology**:
 - Comparing vectors of two nodes provides a more detailed measure of local topological similarity than node degrees or clustering coefficient.



Node Features: Graphlets

Rooted connected induced non-isomorphic subgraphs





There are 30 different graphlets on up to 5 nodes

There are 73 distinct automorphism orbits (or simply orbits) for graphlets with up to 5 nodes.

Orbits are the distinct roles nodes play *within* those graphlets based on symmetry.

Node Features: Graphlets

Rooted connected induced non-isomorphic subgraphs



u has graphlets:


0, 1, 2, 3, 5, 10, 11, ...

Node Features: Graphlets


■ **Graphlet Degree Vector (GDV):** A count vector of graphlets rooted at a given node.

■ **Example:**

Possible graphlets on up to 3 nodes



Graphlet instances of node u :



GDV of node u :

a, b, c, d

$[2, 1, 0, 2]$

Node-Level Features: Summary


- **Importance-based** features: capture the **importance** of a node in a graph
 - **Node degree:**
 - Counts the number of neighboring nodes
 - **Node centrality:**
 - Models the importance of neighboring nodes in a graph
 - Different modeling choices: eigenvector centrality, betweenness centrality, closeness centrality
- Useful for predicting influential nodes in a graph
 - Example: predicting celebrity users in a social network

Node-Level Features: Summary

- **Structure-based** features: capture **topological properties** of the local neighborhood around a node.
 - **Node degree:**
 - Counts the number of neighboring nodes
 - **Clustering coefficient:**
 - Measures how connected neighboring nodes are
 - **Graphlet degree vector:**
 - Counts the occurrences of different graphlets
- Useful for predicting a particular role a node plays in a graph:
 - Example: predicting protein functionality in a protein-protein interaction network.

Link-Level Prediction Task: Recap


- The task is to predict new links based on the existing links.
- At test time, node pairs (with no existing links) are ranked, and top K node pairs are predicted.
- The key is to design features for a pair of nodes.



Link Prediction as a Task


Two formulations of the link prediction task:

- **Links missing at random:**
 - Remove a random set of links and then aim to predict them
- **Links over time:**
 - Given $G[t_0, t'_0]$ a graph defined by edges up to time t'_0 , output a ranked list L of edges (not in $G[t_0, t'_0]$) that are predicted to appear in time $G[t_1, t'_1]$
- **Evaluation:**
 - $n = |E_{new}|$: # new edges that appear during the test period $G[t_1, t'_1]$
 - Take the top n elements of L and count the correct edges




Link Prediction via Proximity

- Methodology:
 - For each pair of nodes (x,y) compute score $\mathbf{c(x,y)}$
 - For example, $c(x,y)$ could be the # of common neighbors of x and y
 - **Sort** pairs (x,y) by the decreasing score $\mathbf{c(x,y)}$
 - **Predict top n** pairs as new links
- See which of these links appear in $G[t_1, t'_1]$




Link-Level Features: Overview

- **Distance-based feature**
- **Local neighborhood overlap**
- **Global neighborhood overlap**



Distance-based Features

- Shortest-path distance between two nodes
- However, this **does not capture the degree of neighborhood overlap**:
 - Node pair (B, H) has 2 shared neighboring nodes, while pairs (B, E) and (A, B) only have 1 such node.
- Example:



$$S_{BH} = S_{BE} = S_{AB} = 2$$

$$S_{BG} = S_{BF} = 3$$

Local Neighborhood Overlap

Captures # neighboring nodes shared between two nodes v_1 and v_2

- **Common neighbors** $|N(v_1) \cap N(v_2)|$

- **Example:**

$$|N(A) \cap N(B)| = |\{C\}| = 1$$

- **Jaccard's coefficient**

- **Example:**

$$\frac{|N(A) \cap N(B)|}{|N(A) \cup N(B)|} = \frac{|\{C\}|}{|\{C, D\}|} = \frac{1}{2}$$


- **Adamic-Adar index**

- **Example:**

$$\frac{1}{\log(k_C)} = \frac{1}{\log 4}$$

$$\frac{|N(v_1) \cap N(v_2)|}{|N(v_1) \cup N(v_2)|}$$

$$\sum_{u \in N(v_1) \cap N(v_2)} \frac{1}{\log(k_u)}$$



Key differences

Common Neighbors:

- Simple absolute count of shared neighbors.
- Higher count = higher similarity, but doesn't consider node size/graph density.

Jaccard Index:


- Normalizes shared count by the total unique neighbors.
- Gives a relative proportion (0 to 1), better for comparing nodes of different sizes.

Adamic-Adar Index:

- Weights shared neighbors by the inverse of their degree.
- More weight given to common neighbors who are less connected themselves (rare).

Global Neighborhood Overlap

- **Limitations** of local neighborhood features:
 - The metric is always zero if the two nodes do not have any neighbors in common.
 - However, the two nodes may still potentially be connected in the future.
- **Global neighborhood overlap metrics resolve the limitation by considering the entire graph.**



$$N_A \cap N_E = \phi$$


$$|N_A \cap N_E| = 0$$

Global Neighborhood Overlap

- **Katz index: count the number of walks of all lengths between a given pair of nodes.**
- **Question:** How to compute #walks between two nodes?
- **Use powers of the graph adjacency matrix!**

Intuition: Powers of Adj Matrices

- Computing #walks between two nodes
 - Recall: $A_{uv} = 1$ if $u \in N(v)$
 - Let $P_{uv}^{(K)} = \text{\#walks of length } K \text{ between } u \text{ and } v$
 - **We will show** $P^{(K)} = A^k$
 - $P_{uv}^{(1)} = \text{\#walks of length 1 (direct neighborhood) between } u \text{ and } v = A_{uv}$



$$A = \begin{pmatrix} 0 & \boxed{1} & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$

$P_{12}^{(1)} = A_{12}$

Intuition: Powers of Adj Matrices

- How to compute $P_{uv}^{(2)}$?
- **Step 1:** compute #walks of length 1 between each of u's neighbors and v
- **Step 2:** Sum up these #walks across u's neighbors

$$P_{uv}^{(2)} = \sum_i A_{ui} * P_{iv}^1 = \sum_i A_{ui} * A_{iv} = A_{uv}^2$$

Node 1's neighbors

#walks of length 1 between
Node 1's neighbors and Node 2

$$A^2 = \begin{pmatrix} 0 & \boxed{1} & 0 & \boxed{1} \\ 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix} \times \begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & \boxed{0} & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & \boxed{1} & 1 & 0 \end{pmatrix} = \begin{pmatrix} 2 & \boxed{1} & 1 & 1 \\ 1 & 2 & 1 & 1 \\ 1 & 1 & 1 & 0 \\ 1 & 1 & 0 & 3 \end{pmatrix}$$

Power of adjacency

Global Neighborhood Overlap

- **Katz index: count the number of walks of all lengths between a pair of nodes.**
- How to compute #walks between two nodes?
- Use adjacency matrix powers!
 - A_{uv} specifies #walks of length 1 (direct neighborhood) between u and v
 - A_{uv}^2 specifies #walks of length 2 between u and v
 - ...
 - A_{uv}^l specifies #walks of length l between u and v

Global Neighborhood Overlap

- Katz index between v_1 and v_2 is calculated as the sum over all walks lengths:

$$S_{v_1 v_2} = \sum_{l=1}^{\infty} \beta^l A_{v_1 v_2}^l$$

$0 < \beta < 1$: attenuation factor (longer paths counts less)

- Katz index matrix is computed in closed form:


$$S = \sum_{i=1}^{\infty} \beta^i A^i = (I - \beta A)^{-1} - I \text{ by geometric series of matrices}$$

Link-Level Features: Summary

- **Distance-based features:**
 - It uses the shortest path length between two nodes but does not capture how the neighborhood overlaps.
- **Local neighborhood overlap:**
 - Captures how many neighboring nodes are shared by two nodes.
 - Becomes zero when no neighbor nodes are shared.
- **Global neighborhood overlap:**
 - It uses a global graph structure to score two nodes.
 - Katz index counts #walks of all lengths between two nodes.

Graph-Level Features

- **Goal**: We want features that characterize the structure of an entire graph.
- For example:



Background: Kernel Methods

- Goal: We want to use machine learning (ML) to make predictions about entire graphs (e.g., classify molecules, detect communities).
- Standard ML Approach: ML models usually work with feature vectors - lists of numbers describing each data point.
- The Problem with Graphs: How do you represent a complex graph as a simple list of numbers (a feature vector)?
- Designing good, informative features by hand that capture the graph's structure can be very difficult and time-consuming.
- **The Kernel Idea:** What if, instead of defining features first, we directly define a way to measure similarity between two graphs? This similarity measure is called a **kernel**.

Graph-Level Features: Overview

Graph Kernels:

- Measure the similarity between two graphs:
 - Graphlet Kernel [1]
 - Weisfeiler-Lehman Kernel [2]
 - Other kernels are also proposed in the literature
 - (beyond the scope of this lecture)
 - Random-walk kernel
 - Shortest-path graph kernel
 - And many more...

1. Shervashidze, Nino, et al. "Efficient graphlet kernels for large graph comparison." Artificial Intelligence and Statistics. 2009.

2. Shervashidze, Nino, et al. "Weisfeiler-lehman graph kernels." Journal of Machine Learning Research 12.9 (2011).

Graph Kernel: Key Idea

- Goal: Design graph feature vector $\Phi(G)$
- Key idea: Bag-of-Words (BoW) for a graph
 - Recall: BoW uses the word counts as document features (no ordering considered).
 - Naïve extension to a graph: regard nodes as words.
 - Since both graphs have 4 red nodes, we get the same feature vector for two different graphs...



Graph Kernel: Key Idea

What if we use **Bag of node degrees**?

Deg1: ● Deg2: ● Deg3: ●

$$\phi(\text{triangle}) = \text{count}(\text{triangle with colored nodes}) = [1, 2, 1]$$

Obtains different features for different graphs!

$$\phi(\text{square}) = \text{count}(\text{square with colored nodes}) = [0, 2, 2]$$

Both Graphlet Kernel and Weisfeiler-Lehman (WL) Kernel use Bag-of-* representation of graph, where * is more sophisticated than node degrees!


Graph-Level Graphlet Features

- **Key idea: Count the number of different graphlets in a graph.**
 - Note: The definition of graphlets here differs slightly from the node-level features.
 - The two differences are:
 - Nodes in graphlets here do not need to be connected (allows for isolated nodes)
 - The graphlets here are not rooted.
 - Examples in the next slide illustrate this.


Graph-Level Graphlet Features

Let $\mathcal{G}_k = (g_1, g_2, \dots, g_{n_k})$ be a list of graphlets of size k .

- For $k = 3$, there are 4 graphlets.



- For $k = 4$, there are 11 graphlets.




Graph-Level Graphlet Features

Given graph G and a graphlet list $G_k = (g_1, g_2, \dots, g_{n_k})$
define the graphlet count vector $f_G \in \mathbb{R}^{n_k}$ as

$$f(G)_i = \#(g_i \subseteq G) \text{ for } i = 1, 2, \dots, n_k$$

Graph-Level Graphlet Features

Example for $k = 3$.



Graph-Level Graphlet Kernel

Given two graphs, G and G' , graphlet kernel is computed as the dot product

$$K(G, G') = f(G)^T f(G')$$

Problem: if G and G' have different sizes, that will greatly skew the value.

Solution: normalize each feature vector

$$h(G) = \frac{f_G}{\text{sum}(f_G)}$$

The Graphlet Kernel

- Limitations: **Counting graphlets is expensive!**
- Counting size- k graphlets for a graph with size n by enumeration takes n^k .
- This is unavoidable in the worst case since the subgraph isomorphism test (judging whether a graph is a subgraph of another graph) is **NP-hard**.
- If a graph's node degree is bounded by d , an $O(nd^{k-1})$ algorithm exists to count all the graphlets of size k .
- **Can we design a more efficient graph kernel?**

Weisfeiler-Lehman Kernel

- **Goal**: Design an efficient graph feature descriptor $\Phi(G)$
- **Idea**: Use neighborhood structure to enrich node vocabulary iteratively.
 - Generalized version of Bag of node degrees since node degrees are one-hop neighborhood information.
- Algorithm to achieve this:
 - **Color refinement**

Color Refinement

Given: A graph G with a set of nodes V .

- Assign an initial color $c^{(0)}(v)$ to each node v .
- Iteratively refine node colors by

$$c^{k+1}(v) = \text{HASH}(c^{(k)}(v), \{c^{(k)}(u)\}_{u \in N(v)})$$



where HASH maps different inputs to different colors

- After k steps of color refinement, $c^{(K)}(v)$ summarizes the structure of the K -hop neighborhood



Color Refinement (1)

Example of color refinement given two graphs

- Assign initial colors





- Aggregate neighboring colors





Color Refinement (2)

Example of color refinement given two graphs

- Aggregated colors



- Hash aggregated colors





Hash table

1,1	-->	2
1,11	-->	3
1,111	-->	4
1,1111	-->	5



Color Refinement (3)

Example of color refinement given two graphs

- Aggregated colors





- Hash aggregated colors




Color Refinement (4)

Example of color refinement given two graphs

- Aggregated colors



- Hash aggregated colors



Hash table

2,4	—>	6
2,5	-->	7
3,44	-->	8
3,45	-->	9
4,245	-->	10
4,345	-->	11
5,2244	-->	12
5,2344	-->	13

Color Refinement (5)

Stopping Condition:


- It stops when **node colors stabilize**.
- Stabilization: **no node's color changes in an iteration based on its current color and neighbors' colors**.

In Practice (for Graph Kernels):

- Often run for a predefined number of iterations (h) for consistent feature vectors.
- Guaranteed to stabilize within the number of nodes.
- Ensure that the resulting feature vectors have the same number of features.
 - NB: the list of unique colors is **global** to all the graphs in the dataset

Weisfeiler-Lehman Graph Features


After color refinement, WL kernel counts number of nodes with a given color.



1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13
= [6, 2, 1, 2, 1, 0, 2, 1, 0, 0, 2, 1, 0]

Colors

Counts



1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13
= [6, 2, 1, 2, 1, 1, 1, 0, 1, 1, 1, 0, 1]

Weisfeiler-Lehman Kernel

The WL kernel value is computed by the inner product of the color count vectors:

$$\begin{aligned} K(\text{graph}_1, \text{graph}_2) &= \phi(\text{graph}_1)^T \phi(\text{graph}_2) \\ &= 49 \end{aligned}$$

Weisfeiler-Lehman Kernel

- **WL kernel is computationally efficient.**
 - The time complexity for color refinement at each step is linear in $\#(\text{edges})$ since it aggregates neighboring colors.
- When computing a kernel value, only colors appearing in the two graphs must be tracked.
 - Thus, at most, $\#(\text{colors})$ is the total number of nodes.
- Counting colors takes linear time w.r.t. $\#(\text{nodes})$.
- In total, **time complexity is linear in $\#(\text{edges})$.**

Graph-level Features: Summary

- **Graphlet Kernel**
 - Graph is represented as Bag-of-graphlets
 - Computationally expensive
- **Weisfeiler-Lehman Kernel**
 - Apply K -step color refinement algorithm to enrich node colors
 - Different colors capture different K -hop neighborhood structures
 - Graph is represented as Bag-of-colors
 - Computationally efficient
 - Closely related to Graph Neural Networks (for the next course!)

Summary

- **Traditional ML Pipeline**
 - Hand-crafted (structural) features + ML model
- **Hand-crafted features for graph data**
 - **Node-level:**
 - Node degree, centrality, clustering coefficient, graphlets
 - **Link-level:**
 - Distance-based feature
 - local/global neighborhood overlap
 - **Graph-level:**
 - Graphlet kernel, WL kernel
- However, we only considered featurizing the graph structure (but not the attribute of nodes and their neighbors)

Q&A 